

AgentRx: A Benchmark Study of LLM Agents for Multimodal Clinical Prediction Tasks

Baraa Al Jorf

New York University Abu Dhabi, UAE

Farah E. Shamout

New York University Abu Dhabi, UAE

BARAA.AL.JORF@NYU.EDU

FS999@NYU.EDU

Abstract

Building effective clinical decision support systems requires the synthesis of complex heterogeneous multimodal data. Such modalities include temporal electronic health records data, medical images, radiology reports, and clinical notes. Large language model (LLM)-based agents have shown impressive performance in various healthcare tasks, especially those involving textual modalities. Considering the fragmentation of healthcare data across hospital systems, collaborative agent frameworks present a promising direction to mitigate data sharing challenges. However, the effectiveness of LLM agents for multimodal clinical risk prediction remains largely unexamined. In this work, we conduct a systematic evaluation of LLM-based agents for clinical prediction tasks using large-scale real-world data. We assess performance in unimodal and multimodal settings and quantify performance gaps between single agent and multi-agent systems. Our findings highlight that single agent frameworks outperform naive multi-agent systems, are better at handling multimodal data, and are better calibrated. This underscores a critical need for improving multi-agent collaboration to better handle heterogeneous inputs. By open-sourcing our code and evaluation framework, this work offers a new benchmark to support future developments relating to agentic systems in healthcare.

Data and Code Availability This study utilizes publicly available data from MIMIC-IV (Johnson et al., 2023b), MIMIC-CXR (Johnson et al., 2019), and MIMIC-IV-Note (Johnson et al., 2023a). Our code and models are open source at: <https://github.com/nyuad-cai/AgentRX>.

Institutional Review Board (IRB) This work did not involve human subjects, so IRB approval was not required.

1. Introduction

The integration of Artificial Intelligence (AI) into clinical decision support systems promises more optimized clinical workflows and better patient outcomes, particularly in resource-constrained settings such as the Intensive Care Unit (ICU). There are many data modalities routinely collected from patients in the ICU, such as Patient Summaries (PS), Electronic Health Records (EHR) data, Chest X-ray (CXR) images, Radiology Reports (RR), and Discharge Notes (DN). Recent work has highlighted the potential of multimodal deep neural networks for fusing such modalities to more accurately predict patient risk outcomes, compared to relying on a single modality (Khader et al., 2023; Lee et al., 2023).

State-Of-The-Art (SOTA) prediction models rely on optimized deep learning-based architectures. For example, MeTra uses a transformer to apply an attention-based fusion mechanism for processing EHR and CXR data (Khader et al., 2023). Similarly, MedFuse uses a simpler Long Short-Term Memory (LSTM) network to fuse latent embeddings of the same modalities (Hayat et al., 2022). MedPatch processes EHR, CXR, RR, and DN data, and uses a modular multistage fusion pipeline (Jorf and Shamout, 2025). While the models achieve high performance metrics in clinical prediction tasks, they have some limitations that impede real-world adoption. First, they are black-box in nature, which undermines clinical trust and interpretability (Catalina et al., 2023; von Eschenbach, 2021; Shuaib, 2024). Second, they typically have fixed data requirements, which restrict

the transferability of these approaches across other modalities and clinical settings (Al Jorf et al., 2026).

Recently, LLMs have demonstrated powerful abilities in emulating human reasoning and text generation (Singhal et al., 2025) and have proven to enhance clinical workflows (Tai-Seale et al., 2024), making them attractive candidates for clinical decision support tools. Recent work specifically highlights their effective ability in generating narrative justifications (Lee et al., 2025). Unlike traditional black box models that rely on abstract feature importance, LLMs can function as semantic translators, converting complex risk predictions into intuitive, natural language explanations that align with clinical reasoning. However, this advantage comes with unreliable performance gains. While some studies highlight the effectiveness of LLMs at handling unimodal EHR or PS data for clinical prediction tasks (Acharya et al., 2024; Kara and Gunel, 2025; Jin et al., 2025), others find that LLMs underperform compared to traditional supervised methods (Tan et al., 2024; Zhu et al., 2025). Notably, these studies have been restricted to unimodal settings. To date, the potential of LLMs to handle multimodal data integration for robust clinical risk prediction remains unexplored.

To address this gap, this work investigates the fundamental question: How effective are agentic systems at multimodal clinical prediction tasks? Our motivation stems from the fragmented nature of healthcare data, where clinical information is often dispersed across multiple isolated databases that ideally would be managed by modality-specific agents to avoid expensive data transfers. However, it remains unclear whether such decentralized, agentic approaches perform well across a variety of data availability settings. Hence, we introduce **AgentRx**, a comprehensive evaluation framework to analyze agent performance across three progressive settings. First, we establish performance baselines using only a single modality, specifically clinical notes encompassed within patient summaries. Second, we assess the capacity of a single agent to synthesize heterogeneous multimodal data within one context window, utilizing modality dropping ablations to measure robustness. Finally, we investigate whether multi-agent reasoning across specialized agents can improve performance compared to single agent approaches.

In summary, we make the following contributions:

- We provide a systemic benchmark (**AgentRx**) for the evaluation of LLM-based clinical prediction

tasks using four data modalities: EHR, RR, CXR, and PS.

- We analyze the impact of data heterogeneity on model performance, in terms of discriminative ability and calibration, by varying the number and type of modalities available to the agents. We also conduct ablations to assess the effect of progressively adding more data modalities in simple single agent and multi-agent setups, to quantify the performance gap between the two.
- We publicly release our code and evaluation framework to enhance the usability of **AgentRx** by the research community and support the advancement of agentic AI in healthcare.

2. Related Work

2.1. LLMs in Healthcare

The application of Natural Language Processing (NLP) in healthcare has rapidly evolved from task-specific discriminative models to general-purpose generative reasoning models. Early approaches relied on encoder-only architectures like BioBERT (Lee et al., 2020) and ClinicalBERT (Huang et al., 2019) for medical tasks like biomedical text mining and hospital readmission prediction. The emergence of LLMs fundamentally shifted this focus towards generative capabilities. Models such as GPT-4 (Bicknell et al., 2024) and Med-PaLM (Singhal et al., 2023) achieved expert-level performance on USMLE-style reasoning and medical question answering. However, while LLMs excel at reasoning tasks like summarizing clinical notes (Afshar et al., 2025; Lukac et al., 2025) and structuring interpretable clinical insights (Lee et al., 2025), their zero-shot ability to forecast temporal clinical outcomes often lags behind traditional supervised baselines (Tan et al., 2024; Zhu et al., 2025). This performance gap suggests that scaling model parameters alone is insufficient for risk stratification, necessitating the development of more structured reasoning frameworks and the integration of multimodal clinical data.

Consequently, Vision-Language Models (VLMs) extend LLM capabilities by aligning visual encoders with language decoders. In the general domain, architectures like Ovis (Lu et al., 2025) and Qwen2.5-VL (Bai et al., 2025) have demonstrated that visual signals can be effectively tokenized and processed

Table 1: **Summary of existing medical multimodal benchmarks for risk prediction.** We summarize the characteristics of existing benchmarks, including both traditional deep learning benchmarks and agentic frameworks. The ‘‘Agentic’’ column indicates whether the benchmark explicitly evaluates agent-based reasoning workflows.

Benchmark	Modalities	Scope	Agentic
DF-Mdl (Chen et al., 2024a)	PS, Medical Images	Clinical Prediction	×
MC-BEC (Chen et al., 2023)	PS, EHR	Clinical Prediction	×
MedPatch (Jorf and Shamout, 2025)	EHR, CXR, RR, DN	Clinical Prediction	×
MedMod (Elsharief et al., 2025)	EHR, CXR	Clinical Prediction	×
EHRXQA (Bae et al., 2023)	EHR, CXR	Medical Question Answering	✓
MedAgents (Tang et al., 2024)	Medical Questions	Medical Question Answering	✓
MDAgents (Kim et al., 2024)	Medical Questions, Images, Videos	Medical Question Answering	✓
AgentRx (Ours)	EHR, CXR, RR, PS	Clinical Prediction	✓

alongside text. Within healthcare, this has led to specialized models such as LLaVA-Med (Li et al., 2023), MedGemma (Sellergren et al., 2025), and HuatuoGPT-Vision (Chen et al., 2024b), which are fine-tuned on biomedical image-text pairs to perform tasks like visual question answering and instruction-following. Similar to their text-only counterparts, medical VLMs are predominantly evaluated on and designed for reasoning tasks rather than predictive ones (Kalpelbe et al., 2025).

2.2. Health Reasoning Agents using LLMs

Since full parameter fine-tuning is often computationally expensive, recent research has pivoted towards optimizing inference-time reasoning to enhance clinical performance. Standard generic techniques have been developed to enhance single agent LLM reasoning. This includes Chain-of-Thought (CoT) prompting (Wei et al., 2022) where the model is allowed to decompose a problem into manageable intermediate steps, Self-Consistency (SC) (Wang et al., 2022) where multiple reasoning paths are used to mitigate errors, Retrieval-Augmented Generation (RAG) where LLMs are provided factual grounding from external knowledge bases (Lewis et al., 2020), and Self-Refinement, which employs iterative critique to correct generated outputs based on feedback (Madaan et al., 2023). These general-purpose strategies have been extensively benchmarked within the medical domain to validate their utility across tasks ranging from medical exams to clinical prediction (Tan et al., 2024).

Building on these foundations, domain-specific frameworks have been designed to further optimize

medical reasoning. A prominent example is Med-prompt (Nori et al., 2023), which synergizes RAG based few-shot selection, CoT reasoning chains, and ensembling to achieve SOTA performance on the MedQA benchmark. Expanding on this, AgentMD (Jin et al., 2025) introduces a tool-learning framework that empowers agents to autonomously curate and apply executable clinical calculators, demonstrating significantly improved accuracy in risk prediction compared to standard LLMs. While effective, these strategies primarily focus on optimizing the output of a single model instance.

2.3. Multi-Agent Systems

Numerous studies have demonstrated that employing multi-agent frameworks can significantly enhance LLM performance on complex reasoning tasks compared to single-agent baselines (Hong et al., 2023; Yang et al., 2025). By simulating human-like collaboration, these systems mitigate individual hallucinations and improve logical consistency. For instance, debate-style frameworks allow agents to critique each other’s outputs until a consensus is reached, effectively filtering out erroneous reasoning steps (Liang et al., 2024; Du et al., 2024). In the medical domain, architectures such as MedAgents (Tang et al., 2024) and MDAgents (Kim et al., 2024) have successfully applied role-playing strategies where agents adopt specific specialist personas to achieve SOTA results on medical question-answering benchmarks. However, preliminary literature suggests that this reasoning advantage may not translate directly to all domains (Gao et al., 2025; Kim et al., 2025; Cemri et al., 2025). For clinical risk prediction, recent benchmarks

indicate that while agents excel at generating explanations, they frequently underperform compared to specialized supervised models in temporal forecasting tasks (Zhu et al., 2025; Tan et al., 2024). Crucially, these negative findings have largely been derived from limited unimodal evaluations, primarily focusing on structured EHR data or text-only inputs. It remains unclear whether the collaborative benefits of multi-agent systems could be better realized in a multimodal setting, where distinct specialized agents could independently process heterogeneous data before engaging in collective decision-making.

2.4. Motivation

We provide an overview of relevant benchmarks in Table 1. Existing benchmarks largely focus on either traditional deep learning models or modern generative reasoning. Most deep learning based benchmarks focus on clinical prediction tasks, such as in-hospital mortality, length of stay, or patient phenotyping (Bae et al., 2023; Elsharief et al., 2025), while agentic benchmarks focused mostly on reasoning and retrieval capabilities for medical question and answering. For example, MedAgents and MDAgents have established rigorous standards for multimodal medical LLMs, but predominantly focused on Medical Question Answering (Tang et al., 2024; Kim et al., 2024). While current LLM medical benchmarks evaluate an agent’s ability to retrieve knowledge or answer complex queries, they do not assess the agent’s capacity for temporal forecasting or risk stratification in a clinical multimodal setting. The limited agentic prediction benchmarks that do exist are strictly unimodal (Tan et al., 2024; Zhu et al., 2025). This distinction highlights a critical need for benchmarks that specifically evaluate VLMs on predictive clinical endpoints, which is the main focus of our study.

3. Methodology

3.1. Preliminaries

To formalize the **AgentRx** benchmark, we introduce relevant notation. For a given patient encounter p , we assume the presence of multimodal data $\mathcal{X}_p = \{\mathbf{x}_{ps}, \mathbf{x}_{ehr}, \mathbf{x}_{cxr}, \mathbf{x}_{rr}\}$. Here, \mathbf{x}_{ps} represents a textual patient summary which includes the patient age, sex, and clinical history, $\mathbf{x}_{ehr} \in \mathbb{R}^{d \times t}$ represents the multivariate time-series data from the patient’s EHR with d features over t time steps. $\mathbf{x}_{cxr} \in \mathbb{R}^{h \times w}$ represents a CXR image with height and width $h, w = 224$ and \mathbf{x}_{rr}

Table 2: **Dataset statistics per modality and task.** We describe the sample counts for the training and test splits across the mortality and length of stay prediction tasks. The counts are stratified by modality.

Modality	Mortality		Length of Stay	
	Training	Test	Training	Test
PS	17,773	4,925	17,476	4,845
EHR	17,773	4,925	17,476	4,845
RR	16,128	4,454	15,865	4,380
CXR	4,259	1,174	4,171	1,153

represents the textual radiology reports associated with all medical images collected from the patient. The goal of the system is to predict a groundtruth label denoted by y .

3.2. Real-world Multimodal Dataset

3.2.1. DATA CURATION

We extracted the EHR data from MIMIC-IV (Johnson et al., 2023b), the CXR images from MIMIC-CXR (Johnson et al., 2019), and the RR and PS from MIMIC-IV-Notes (Johnson et al., 2023a). The MIMIC-IV dataset includes de-identified data from over 315,460 patients with ICU stays at the Beth Israel Deaconess Medical Center between 2008 and 2019. The MIMIC-CXR dataset consists of over 377,000 chest radiographs. MIMIC-IV-Note supplements these two datasets with unstructured textual data, including both 331,794 DNs and 2,321,355 RRs. We constructed the multimodal dataset by aligning the subject, stay, and admission identifiers across all three datasets. We mandate the existence of PS for all samples, while other modalities are paired if available.

3.2.2. PREDICTION TASKS

We introduce two clinical prediction tasks and perform evaluations across two modality settings, with a unified evaluation scheme. We report the Area Under the Receiver Operating Characteristic Curve (AUROC), Area Under the Precision-Recall Curve (AUPRC), and Expected Calibration Error (ECE):

- In-hospital mortality prediction is a binary classification task that involves predicting risk of in-hospital mortality based on the first 48 hours in the ICU.

- Long Length of Stay prediction is a binary classification task that involves predicting whether a patient’s stay is going to extend beyond 7 days, by the end of the first 48 hours in the ICU.

After multimodal pairing, the dataset was split into training (70%), validation (10%), and testing (20%) sets. The exact dataset distribution per task is reported in Table 2.

3.2.3. MULTIMODAL DATA PROCESSING AND PAIRING

Since patient summaries are considered to be the primary base modality, we ensure that every patient in our cohort has one. We generated patient summaries by processing the raw DNs from MIMIC-IV-Note. To prevent data leakage, we extracted only the information available prior to ICU admission. This includes the patient’s medical, surgical, and family histories as well as the demographic details such as age and sex, ensuring the agent operates with the same context available to a clinician at the time of admission.

For EHR, we used a set of 17 clinical variables consistent with previous work (Hayat et al., 2022; Jorf and Shamout, 2025). These include 5 categorical variables (capillary refill rate, Glasgow coma scale eye opening, motor response, verbal response, and total) and 12 continuous variables (diastolic blood pressure, fraction of inspired oxygen, glucose, heart rate, height, mean blood pressure, oxygen saturation, respiratory rate, systolic blood pressure, temperature, weight, and pH). To format the EHR data for the LLM agents, we serialized it into a structured text format $[T_0+\Delta T]$ Variable=Value where T_0 is the time of admission, recording only observed measurements at each timestamp to minimize token usage and handle irregular sampling rates efficiently. The observation window was strictly limited to the first 48 hours of ICU admission. To strictly adhere to context window limits, high-frequency streams exceeding 500 time-steps were truncated to preserve the initial admission state (first 100 steps) and the most recent clinical trajectory (last 400 steps).

For the same cohort, we included CXR images that were collected within the first 48 hours of ICU admission. We restricted the selection to Anterior-Posterior (AP) views, as these are standard for portable ICU bedside imaging. For patients with multiple scans within the window, we selected the latest valid scan to capture the most recent clinical state prior to the prediction horizon.

We extracted reports corresponding to patients available in the dataset. MIMIC-IV-Note includes reports for various imaging modalities (CT, MRI, Ultrasound). We concatenated all relevant reports for each patient that were collected within 48 hours from admission into a single report aggregate that is passed as the full RR modality.

3.3. Agentic Frameworks

We formalize the inference logic for our agentic evaluation framework. We define an agent $\mathcal{A}(\cdot)$ as a functional unit that processes clinical context to produce a prediction or a probability. We define \mathcal{P}_{task} as the system prompt directing the agent’s reasoning for unimodal and multimodal tasks, and explore three main settings depicted in Figure 1:

- Setting 1: Single Agent Unimodal
In this setting, the agent relies solely on the Patient Summary (\mathbf{x}_{ps}) to form a baseline prediction as demonstrated in Appendix A1.
- Setting 2: Single Agent Multimodal
Here, a single generalist agent processes all available modalities $\mathcal{X}_{available}$ simultaneously within a single context window as demonstrated in Appendix A2.
- Setting 3: Multi-Agent Multimodal
This setting employs specialized agents \mathcal{A}_m for each modality m . Each agent independently generates a probability p_m , which is then averaged. This is shown in Appendix A3

3.4. Baselines

To assess the efficacy of different agentic architectures, we compare performance across the following single-agent and multi-agent baselines.

3.4.1. SUPERVISED BASELINES

We compare our agentic frameworks against two specialized deep learning architectures:

1. BioBERT (Unimodal): For the text-only setting, we utilize BioBERT (Lee et al., 2020), a language model pre-trained on biomedical corpora (PubMed). We freeze the backbone and fine-tune a linear classification head on the token embeddings of the PS to establish a strong supervised baseline.

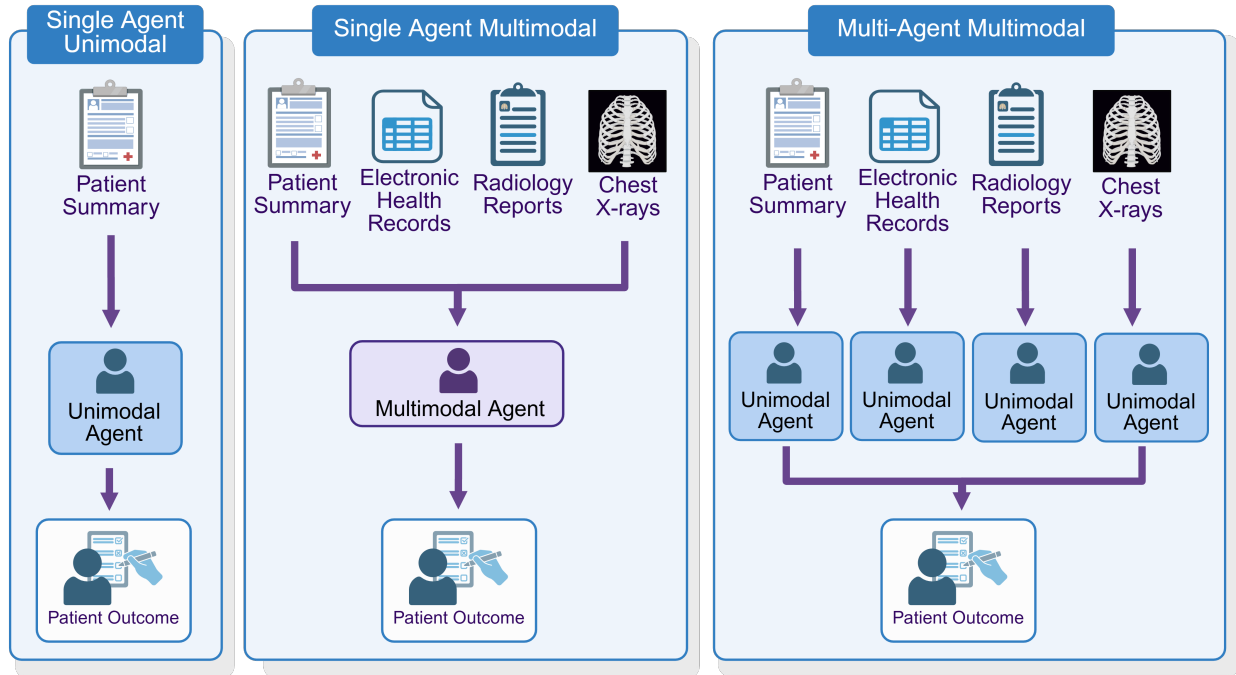


Figure 1: Overview of the agentic evaluation frameworks considered within the **AgentRx** benchmark, spanning Single-Agent (Unimodal/Multimodal) and Multi-Agent settings.

2. MedPatch (Multimodal): For the multimodal setting, we employ MedPatch (Jorf and Shamout, 2025), a SOTA fusion architecture that utilizes a confidence-guided patching mechanism to effectively integrate heterogeneous modalities.

3.4.2. SINGLE-AGENT BASELINES - UNIMODAL AND MULTIMODAL

1. Zero-shot: Vanilla baseline where we feed the model all the setting’s available modalities and ask it for a prediction.
2. Few-shot: A baseline where we feed the model one positive example and one negative example (data + labels from the training set) and then ask it to predict the outcome for a test sample (Brown et al., 2020).
3. Chain-of-Thought (CoT): A baseline where we allow the model to first generate a reasoning step, then use that reasoning to produce a prediction (Wei et al., 2022).
4. Self-Consistency + Chain-of-Thought (CoT-SC): A baseline where the model runs 3 parallel

reasoning pathways and then makes a decision using all three paths via voting (probability averaging) (Wang et al., 2022).

5. Self-Refinement: A baseline where the model generates a prediction with reasoning, then self-evaluates its reasoning before making a final prediction based on the evaluation feedback (Madaan et al., 2023).

3.4.3. MULTI-AGENT BASELINES

1. Majority Vote: A baseline where unimodal agents independently analyze their specific data and vote on the outcome (Kaesberg et al., 2025).
2. Debate: A baseline where unimodal agents debate with each other until they reach a consensus (Du et al., 2024).
3. Meta-Prompting: A baseline where a meta-agent evaluates the data and either makes a prediction or instantiates other expert agents to help refine its task (Hou et al., 2022).
4. Traj-CoA + Multimodal Judge: A baseline that uses multiple worker agents to construct an EHR

memory. Each worker receives a chunk of EHR data. The EHR memory and the other modalities are then passed to a multimodal judge agent that makes a final prediction. This baseline combines a unimodal and a multimodal agent setup (Zeng et al., 2025).

5. MDAgents: A baseline where a diverse ensemble of agents is initialized depending on patient case severity (Kim et al., 2024).
6. MedAgents: A baseline where agents follow pre-defined roles and collaborate to create a patient state summary to enable outcome prediction (Tang et al., 2024).

3.5. Implementation Details

To ensure a rigorous evaluation, we instantiate all agentic baselines using four distinct VLM backbones, two general-purpose and two specialized medical. All models utilize approximately 7-8 billion parameters to maintain comparable computational requirements.

3.5.1. MODEL BACKBONES

1. **Qwen2.5-VL-7B-Instruct** (Bai et al., 2025): A leading generalist VLM built upon the Qwen2.5 language model. It features a specialized vision encoder optimized for high-resolution image understanding, making it a strong generalist baseline.
2. **InternVL2.5-8B-MPO** (Chen et al., 2025): A general-purpose VLM from OpenGVLab designed for robust visual reasoning. This variant employs Mixed Preference Optimization (MPO) to align the model’s outputs with human preference.
3. **HuatuoGPT-Vision-7B-Qwen2.5VL** (Chen et al., 2024b): A specialized medical VLM initialized from Qwen2.5-VL. It is further post-trained on a massive corpus of multimodal medical data.
4. **LLaVA-Med-v1.5-Mistral-7B** (Li et al., 2023): A biomedical VLM that builds on the Mistral-7B language model. It was pre-trained on a large-scale dataset of biomedical figure-caption pairs (PMC-15M).

3.5.2. HYPERPARAMETERS

All experiments utilized a standard batch size of 3 or 4, with exceptions for the supervised architectures (batch size 16) and multimodal LLaVA-Med, which required a batch size of 1 to accommodate heterogeneous modality combinations. Regarding context constraints, the high token overhead of serialized tabular EHR data exceeded the effective limit of InternVL2.5, necessitating its evaluation on a restricted 3-modality subset (PS, CXR, RR) in Appendix B1.

4. Results

4.1. Unimodal Results

In Table 3, we present the Single Agent Unimodal setup where agents are restricted to the PS modality. For the mortality prediction task, specialized LLMs demonstrate the capacity to exceed supervised baselines. While BioBERT achieves a competitive AUROC of 0.680, the medical-specific backbone HuaTuo surpasses this in the few-shot setting with an AUROC of 0.700. In terms of precision-recall balance, HuaTuo Zero-shot attains the highest AUPRC of 0.238, improving upon the BioBERT baseline of 0.228. Generalist models also show competence with Qwen few-shot achieving an AUROC of 0.697, closely following the specialized medical agents.

Despite strong discriminative performance (AUROC/AUPRC), generative backbones exhibit significant miscalibration compared to discriminative models. The supervised BioBERT baseline maintains a minimal ECE of 0.006. Conversely, all agentic setups yield considerably higher ECE values. For instance, the top-performing HuaTuo Few-shot achieves an ECE of 0.093, while Llava configurations consistently exceed 0.750. This indicates that while agents can rank patient risk effectively, their probabilistic confidence scores are less reliable than the supervised baseline.

For LoS prediction, BioBERT achieves the highest performance across all metrics with an AUROC of 0.641 and AUPRC of 0.307. Among the agentic backbones, the generalist Intern model performs best, with the zero-shot setup achieving an AUROC of 0.620 and AUPRC of 0.278. Notably, the medical-specific HuaTuo backbone underperforms in this task (zero-shot AUROC 0.599), suggesting that clinical pre-training may be more beneficial for diagnostic tasks like mortality than for operational metrics like LOS.

Table 3: **Single Agent Unimodal Results.** Comparison of agentic backbones against a supervised baseline (BioBERT) using only the PS. Metrics are reported with 95% Confidence Intervals. The best overall performance in each column is bolded.

Backbone	Method	In-Hospital Mortality			Length of Stay (>7 Days)		
		AUROC	AUPRC	ECE	AUROC	AUPRC	ECE
Supervised	BioBERT	0.680 (0.657 - 0.702)	0.228 (0.203 - 0.263)	0.006	0.641 (0.621 - 0.661)	0.307 (0.282 - 0.337)	0.024
Qwen	Zero-shot	0.667 (0.645 - 0.690)	0.234 (0.208 - 0.267)	0.025	0.603 (0.582 - 0.624)	0.261 (0.241 - 0.286)	0.105
	Few-shot	0.697 (0.678 - 0.718)	0.236 (0.213 - 0.270)	0.042	0.602 (0.580 - 0.621)	0.257 (0.237 - 0.280)	0.122
	CoT	0.650 (0.629 - 0.675)	0.214 (0.191 - 0.244)	0.060	0.592 (0.573 - 0.613)	0.260 (0.240 - 0.286)	0.480
	CoT-SC	0.679 (0.657 - 0.702)	0.236 (0.209 - 0.269)	0.029	0.593 (0.571 - 0.613)	0.259 (0.239 - 0.284)	0.454
	Self-Refine	0.666 (0.644 - 0.689)	0.212 (0.187 - 0.240)	0.080	0.580 (0.560 - 0.601)	0.236 (0.219 - 0.257)	0.198
Intern	Zero-shot	0.672 (0.650 - 0.695)	0.216 (0.193 - 0.245)	0.090	0.620 (0.599 - 0.641)	0.278 (0.256 - 0.304)	0.125
	Few-shot	0.687 (0.665 - 0.709)	0.223 (0.199 - 0.251)	0.074	0.605 (0.585 - 0.626)	0.261 (0.241 - 0.286)	0.183
	CoT	0.640 (0.617 - 0.666)	0.202 (0.180 - 0.228)	0.059	0.568 (0.548 - 0.587)	0.224 (0.208 - 0.244)	0.220
	CoT-SC	0.659 (0.637 - 0.683)	0.221 (0.197 - 0.249)	0.086	0.592 (0.570 - 0.612)	0.253 (0.234 - 0.279)	0.239
	Self-Refine	0.630 (0.606 - 0.654)	0.189 (0.169 - 0.216)	0.065	0.568 (0.548 - 0.586)	0.224 (0.207 - 0.244)	0.151
HuaTuo	Zero-shot	0.692 (0.672 - 0.714)	0.238 (0.213 - 0.268)	0.166	0.599 (0.578 - 0.620)	0.256 (0.236 - 0.281)	0.203
	Few-shot	0.700 (0.681 - 0.721)	0.227 (0.205 - 0.258)	0.093	0.602 (0.580 - 0.621)	0.260 (0.240 - 0.284)	0.029
	CoT	0.670 (0.650 - 0.693)	0.219 (0.194 - 0.247)	0.211	0.595 (0.574 - 0.615)	0.258 (0.239 - 0.280)	0.599
	CoT-SC	0.670 (0.649 - 0.692)	0.218 (0.194 - 0.246)	0.211	0.594 (0.574 - 0.614)	0.258 (0.239 - 0.281)	0.599
	Self-Refine	0.671 (0.651 - 0.694)	0.213 (0.190 - 0.242)	0.106	0.581 (0.560 - 0.601)	0.249 (0.230 - 0.273)	0.385
Llava	Zero-shot	0.642 (0.621 - 0.666)	0.204 (0.182 - 0.235)	0.755	0.574 (0.552 - 0.594)	0.234 (0.216 - 0.255)	0.787
	Few-shot	0.684 (0.664 - 0.706)	0.233 (0.208 - 0.265)	0.776	0.605 (0.584 - 0.624)	0.253 (0.235 - 0.273)	0.780
	CoT	0.633 (0.610 - 0.655)	0.170 (0.154 - 0.186)	0.831	0.509 (0.499 - 0.518)	0.194 (0.183 - 0.205)	0.808
	CoT-SC	0.645 (0.622 - 0.666)	0.197 (0.174 - 0.221)	0.840	0.542 (0.523 - 0.560)	0.206 (0.194 - 0.221)	0.806
	Self-Refine	0.631 (0.608 - 0.655)	0.170 (0.154 - 0.187)	0.829	0.516 (0.504 - 0.526)	0.196 (0.185 - 0.208)	0.808

4.2. Multimodal Results

Table 4 presents the performance of agentic backbones when processing multimodal clinical data. A performance gap persists between generalist agents and the supervised MedPatch for both tasks. MedPatch achieves an AUROC of 0.877 for in-hospital mortality and 0.844 for length of stay, whereas the best-performing agent configurations trail significantly behind at 0.763 (HuaTuo Few-shot) and 0.714 (Qwen Zero-shot). This disparity highlights a structural limitation: while LLMs possess strong semantic reasoning capabilities, they lack the dedicated fusion layers found in specialized architectures optimized for multimodal inputs. Despite the gap however, integrating multiple modalities still yields considerable improvements over unimodal baselines in the single agent settings. For the Qwen backbone, mortality prediction performance improves from a unimodal baseline of 0.667 to 0.756 in the multimodal zero-shot setting. Likewise, LOS performance for Qwen increases from 0.603 to 0.714.

We test the hypothesis that distributing reasoning across specialized agents improves outcomes. Collaborative protocols often result in performance degradation. In Table 4, Qwen-based Debate and Meta-Prompt architectures yield mortality AUROCs of

0.631 and 0.599, respectively, lower than the single agent zero-shot baseline (0.756). Among multi-agent frameworks, Traj-CoA consistently performs best (e.g., Qwen Mortality 0.762). Unlike fully decentralized baselines, this architecture employs a final decision-maker that retains direct access to raw multimodal data while incorporating the intermediate reasoning trajectories of the specialized EHR agents.

The supervised MedPatch baseline maintains a low ECE of 0.0189. While the Qwen single agent zero-shot setup remains relatively calibrated (ECE 0.0233), the majority vote architecture exhibits an increase in calibration error to 0.111. This indicates that voting mechanisms may distort the probabilistic confidence of the system. Intern Traj-CoA defies this trend in specific instances, achieving an ECE of 0.010 for mortality prediction. Another observation is that different backbone have varying capabilities when it comes to handling multimodal data. This is evident in Appendix B2 where different backbones show different reductions in ECE with added modalities. Some like Intern show an increase in ECE with more modalities added.

Table 4: **Multimodal Results.** Comparison of Supervised, Single-Agent, and Multi-Agent architectures on the 4-modality dataset. Best overall performance in each column is bolded.

Backbone	Arch.	Method	In-Hospital Mortality			Length of Stay (>7 Days)		
			AUROC	AUPRC	ECE	AUROC	AUPRC	ECE
Supervised	-	MedPatch	0.877 (0.864 - 0.888)	0.546 (0.504 - 0.585)	0.019	0.844 (0.830 - 0.857)	0.551 (0.517 - 0.587)	0.025
Qwen	Single	Zero-shot	0.756 (0.737 - 0.776)	0.330 (0.297 - 0.368)	0.023	0.714 (0.694 - 0.731)	0.345 (0.320 - 0.373)	0.411
		Few-shot	0.763 (0.744 - 0.783)	0.325 (0.292 - 0.361)	0.032	0.682 (0.662 - 0.699)	0.318 (0.294 - 0.346)	0.479
		CoT	0.733 (0.713 - 0.752)	0.274 (0.244 - 0.308)	0.049	0.683 (0.663 - 0.702)	0.311 (0.288 - 0.336)	0.676
		CoT-SC	0.762 (0.742 - 0.782)	0.337 (0.300 - 0.379)	0.039	0.698 (0.678 - 0.715)	0.340 (0.313 - 0.370)	0.661
	Multi	Majority Vote	0.748 (0.727 - 0.768)	0.315 (0.282 - 0.355)	0.111	0.710 (0.690 - 0.728)	0.352 (0.324 - 0.383)	0.046
		Debate	0.631 (0.607 - 0.656)	0.210 (0.185 - 0.243)	0.091	0.644 (0.623 - 0.664)	0.281 (0.259 - 0.307)	0.121
		Meta-Prompt	0.599 (0.573 - 0.625)	0.179 (0.160 - 0.204)	0.051	0.537 (0.514 - 0.558)	0.226 (0.208 - 0.249)	0.086
		Traj-CoA	0.762 (0.743 - 0.780)	0.318 (0.283 - 0.355)	0.039	0.708 (0.688 - 0.726)	0.336 (0.310 - 0.365)	0.190
		MDAgents	0.624 (0.601 - 0.647)	0.192 (0.169 - 0.221)	0.110	0.584 (0.563 - 0.603)	0.240 (0.222 - 0.262)	0.138
		MedAgents	0.662 (0.641 - 0.686)	0.206 (0.185 - 0.235)	0.034	0.634 (0.613 - 0.654)	0.285 (0.261 - 0.310)	0.019
HuaTuo	Single	Zero-shot	0.762 (0.743 - 0.782)	0.325 (0.296 - 0.362)	0.049	0.704 (0.686 - 0.721)	0.331 (0.307 - 0.358)	0.418
		Few-shot	0.763 (0.744 - 0.782)	0.324 (0.291 - 0.361)	0.167	0.703 (0.684 - 0.720)	0.332 (0.308 - 0.358)	0.514
		CoT	0.697 (0.678 - 0.720)	0.233 (0.208 - 0.266)	0.306	0.639 (0.617 - 0.658)	0.282 (0.260 - 0.306)	0.671
		CoT-SC	0.696 (0.678 - 0.719)	0.233 (0.208 - 0.265)	0.305	0.639 (0.617 - 0.658)	0.282 (0.261 - 0.306)	0.669
	Multi	Majority Vote	0.711 (0.690 - 0.730)	0.245 (0.217 - 0.279)	0.050	0.691 (0.671 - 0.709)	0.322 (0.296 - 0.352)	0.087
		Debate	0.628 (0.604 - 0.652)	0.188 (0.165 - 0.216)	0.109	0.621 (0.599 - 0.642)	0.273 (0.250 - 0.300)	0.061
		Meta-Prompt	0.636 (0.614 - 0.657)	0.179 (0.160 - 0.202)	0.428	0.558 (0.534 - 0.578)	0.228 (0.210 - 0.250)	0.080
		Traj-CoA	0.744 (0.725 - 0.766)	0.295 (0.266 - 0.333)	0.065	0.701 (0.680 - 0.718)	0.326 (0.301 - 0.353)	0.323
		MDAgents	0.490 (0.467 - 0.514)	0.125 (0.112 - 0.142)	0.309	0.571 (0.551 - 0.590)	0.237 (0.218 - 0.260)	0.059
		MedAgents	0.606 (0.584 - 0.629)	0.165 (0.148 - 0.187)	0.335	0.584 (0.563 - 0.606)	0.230 (0.213 - 0.252)	0.486
Llava	Single	Zero-shot	0.741 (0.721 - 0.761)	0.268 (0.242 - 0.303)	0.835	0.613 (0.596 - 0.630)	0.242 (0.226 - 0.259)	0.803
		Few-shot	0.684 (0.662 - 0.706)	0.225 (0.199 - 0.255)	0.704	0.581 (0.561 - 0.602)	0.233 (0.217 - 0.252)	0.786
		CoT	0.676 (0.658 - 0.693)	0.184 (0.167 - 0.202)	0.819	0.553 (0.537 - 0.567)	0.210 (0.197 - 0.223)	0.807
		CoT-SC	0.691 (0.670 - 0.714)	0.215 (0.192 - 0.240)	0.806	0.583 (0.563 - 0.602)	0.226 (0.211 - 0.242)	0.806
	Multi	Majority Vote	0.710 (0.686 - 0.732)	0.280 (0.248 - 0.316)	0.812	0.674 (0.654 - 0.693)	0.291 (0.268 - 0.316)	0.765
		Debate	0.495 (0.470 - 0.519)	0.121 (0.109 - 0.137)	0.293	0.506 (0.485 - 0.525)	0.192 (0.178 - 0.210)	0.437
		Meta-Prompt	0.659 (0.638 - 0.683)	0.211 (0.189 - 0.237)	0.810	0.576 (0.555 - 0.594)	0.231 (0.214 - 0.250)	0.780
		Traj-CoA	0.694 (0.673 - 0.718)	0.249 (0.221 - 0.283)	0.758	0.608 (0.589 - 0.626)	0.239 (0.223 - 0.257)	0.795
		MDAgents	0.565 (0.544 - 0.588)	0.143 (0.130 - 0.160)	0.771	0.492 (0.473 - 0.512)	0.194 (0.181 - 0.210)	0.759
		MedAgents	0.609 (0.589 - 0.633)	0.164 (0.147 - 0.184)	0.683	0.545 (0.525 - 0.563)	0.207 (0.193 - 0.223)	0.782

4.3. Ablations

Table 5 presents an ablation study for the in-hospital mortality task, revealing that single agent architectures benefit significantly from increased modality integration. As predictive performance scales positively with the addition of modalities, calibration also improves, with the Qwen backbone seeing AUROC rise from 0.667 to 0.756 while its ECE drops to 0.0233. A similar trend appears for HuaTuo and LLava. This supports the idea that single agents effectively leverage cross-modal dependencies within heterogeneous data to enhance performance.

In contrast, majority vote demonstrates a divergence in reliability, exhibiting degrading calibration as the system complexity increases. Despite the contribution of specialized agents to classification accuracy, the aggregation of hard labels without shared context leads to significant system overconfidence, illustrated by Qwen’s ECE rising from 0.068 to 0.111 in the full multimodal setting. This suggests that decentralized voting mechanisms fail to capture the un-

certainty reduction benefits inherent in multimodal data, unlike the unified processing found in single-agent setups. The ablation results are visualized in Appendix C1. Considering the different performance improvements noticed as modalities are added to the majority vote, we experiment with a weighted majority vote baseline in Appendix C2. In that baseline, we weigh each modality agent by its unimodal AUROC and notice minor performance improvements.

To further understand the failure modes of multi-agent systems, we analyze the generated traces and conduct more ablations. Table 6 shows the percentage of samples in the debate baseline traces that reached consensus at each round, or never reached consensus and had their probabilities averaged (MAX rounds). The results show a big variety depending on the backbone used. For instance, the Qwen backbone reached consensus for all the samples from the first round (indicating no inter-agent debate), while LLavaMed never reached a consensus for half of the patients. Notably, the models show a trend where

Table 5: **Modality Ablations.** Comparing zero-shot vs. majority vote settings for in-hospital mortality. Note that while AUROC/AUPRC scale with modality density, Single Agent calibration (ECE) improves with more data while majority vote calibration deteriorates. Best overall performance in each column is bolded.

Backbone	Arch.	Modalities	In-Hospital Mortality		
			AUROC	AUPRC	ECE
Qwen	Single (ZS)	PS	0.667 (0.645 - 0.690)	0.234 (0.208 - 0.267)	0.025
		PS + CXR	0.671 (0.649 - 0.695)	0.236 (0.209 - 0.267)	0.027
		PS + CXR + RR	0.732 (0.712 - 0.753)	0.273 (0.244 - 0.305)	0.027
		All (PS+EHR+CXR+RR)	0.756 (0.737 - 0.776)	0.330 (0.297 - 0.368)	0.023
	Multi (MV)	PS + CXR	0.666 (0.643 - 0.689)	0.221 (0.197 - 0.252)	0.068
		PS + CXR + RR	0.725 (0.704 - 0.745)	0.266 (0.237 - 0.301)	0.089
All (PS+EHR+CXR+RR)		0.748 (0.727 - 0.768)	0.315 (0.282 - 0.355)	0.111	
HuaTuo	Single (ZS)	PS	0.692 (0.672 - 0.714)	0.238 (0.213 - 0.268)	0.166
		PS + CXR	0.695 (0.675 - 0.716)	0.236 (0.212 - 0.266)	0.205
		PS + CXR + RR	0.742 (0.723 - 0.763)	0.277 (0.248 - 0.310)	0.269
		All (PS+EHR+CXR+RR)	0.762 (0.743 - 0.782)	0.325 (0.296 - 0.362)	0.049
	Multi (MV)	PS + CXR	0.649 (0.629 - 0.670)	0.201 (0.177 - 0.229)	0.068
		PS + CXR + RR	0.689 (0.667 - 0.709)	0.220 (0.195 - 0.252)	0.014
All (PS+EHR+CXR+RR)		0.711 (0.690 - 0.730)	0.245 (0.217 - 0.279)	0.050	
Llava	Single (ZS)	PS	0.642 (0.621 - 0.666)	0.204 (0.182 - 0.235)	0.755
		PS + CXR	0.658 (0.636 - 0.681)	0.213 (0.188 - 0.243)	0.766
		PS + CXR + RR	0.734 (0.712 - 0.756)	0.270 (0.241 - 0.304)	0.779
		All (PS+EHR+CXR+RR)	0.741 (0.721 - 0.761)	0.268 (0.242 - 0.303)	0.835
	Multi (MV)	PS + CXR	0.621 (0.597 - 0.647)	0.184 (0.164 - 0.211)	0.804
		PS + CXR + RR	0.694 (0.672 - 0.717)	0.256 (0.227 - 0.291)	0.800
All (PS+EHR+CXR+RR)		0.711 (0.690 - 0.730)	0.245 (0.217 - 0.279)	0.812	

Table 6: Consensus analysis showing the percentage of samples reaching consensus at each round and AUROC.

Backbone	Round 1	Round 2	Round 3	MAX	AUROC
LlaVaMed	47.1%	0.2%	0.2%	52.6%	0.495
Huatuo	96.8%	2.1%	0.3%	0.8%	0.628
Qwen	100%	0%	0%	0%	0.631

having more samples with more rounds of debate correlates with degraded performance in AUROC. This suggests that current debate capabilities are lacking and do not consistently improve agentic reasoning. Beyond consensus, Appendix C3 summarizes the echo chamber behavior observed across the backbones when used for the debate baseline. We also experiment with a debate between multimodal agents (in Appendix C4), where 4 full-modality agents undergo 3 rounds of debate until consensus. We note that the

performance there still lags behind the single agent baselines. Considering the trends in consensus, sycophancy, and the multimodal debate results, it seems that multi-agent systems suffer from a combination of information loss, noise propagation, and weak modality fusion. This is further evidenced by the Traj-CoA baseline’s performance, where the presence of a multimodal decision-making agent, closely resembling the single-agent setup, reduces the observed errors.

5. Discussion

In this work, we introduced **AgentRx**, a first-of-its-kind benchmark for evaluating LLM-based agentic systems on high-stakes clinical risk prediction tasks. We systematically analyzed how different agentic architectures, ranging from single-agent baselines to complex multi-agent frameworks, perform against supervised deep learning models. The study yields a critical insight that multi-agent multimodal systems

consistently underperform single-agent multimodal systems. This is driven by a fundamental divergence in system calibration where we observe a striking contrast in how single agent systems improve calibration with additional data modalities while multi-agent systems degrade. Closing this calibration gap is vital for the development of scalable decentralized healthcare systems.

More generally, our findings align with the recent benchmarks established by [Zhu et al. \(2025\)](#) in **MedAgentBoard**, confirming that agentic systems generally trail state-of-the-art supervised fusion networks in clinical prediction. However, we identify a critical modality-dependent nuance to this observation. The analysis in **MedAgentBoard** largely focuses on structured EHR data, where LLMs struggle to encode high-dimensional numerical features. In contrast, our unimodal setting utilizing PS which is a free-text modality demonstrates that specialized agents can effectively outperform supervised baselines. This suggests that the performance drop is not intrinsic to the task of risk prediction, but rather a consequence of the modality. Notably, LLM agents seem to synthesize multimodal data more effectively when the task at hand is diagnostic in nature. This is shown in [Appendix C5](#), where both the single agent and multi-agent systems outperform the supervised baselines in the multimodal setting on the note-based chronic kidney disease detection task.

Despite these contributions, this study has limitations. Our benchmark relies on the MIMIC database, which represents a single-center cohort and may pose challenges for generalizability. Additionally, we have limited our evaluation to architectures of sizes within 7-8 billion parameters. We acknowledge that smaller or larger models may exhibit different reasoning capabilities. [Appendix C6](#) summarizes the performance of a smaller backbone family on some frameworks from our benchmark. Finally, the serialization of high-frequency EHR data into text-based context windows remains an open research problem that likely limits the agents' ability to capture subtle physiological trends. In [Appendix C7](#), we explore some alternative serialization strategies. Moving forward, we plan to expand **AgentRx** to include more diverse clinical endpoints and investigate architectures that combine the interpretability of LLM reasoning with the predictive precision of frozen, supervised encoders. Given the current limitations of the examined baselines in handling complex multimodal data, future work could focus on moving away from purely textual communi-

cation and exploring latent-based representations or ICU-specific tool-calling frameworks to enhance reasoning and risk prediction.

Acknowledgements

This work was supported by ASPIRE, the technology program management pillar of Abu Dhabi's Advanced Technology Research Council (ATRC), via the ASPIRE Precision Medicine Research Institute Abu Dhabi (ASPIREPMRIAD) award grant number VRI-20-10, the NYUAD Center for Artificial Intelligence and Robotics, funded by Tamkeen under the NYUAD Research Institute Award CG010, the Meem Foundation, and the ADIA Lab Health Sciences Grant. The research was carried out on the High Performance Computing resources at New York University Abu Dhabi. [Figure 1](#) was created in BioRender (Al Jorf, B. (2025) <https://BioRender.com/jl2gr7s>).

References

- Angeela Acharya, Sulabh Shrestha, Anyi Chen, Joseph Conte, Sanja Avramovic, Siddhartha Sikdar, Antonios Anastasopoulos, and Sanmay Das. Clinical risk prediction using language models: benefits and considerations. *Journal of the American Medical Informatics Association: JAMIA*, 31(9):1856–1864, September 2024. ISSN 1527-974X. doi: 10.1093/jamia/ocae030.
- Majid Afshar, Mary Ryan Baumann, Felice Resnik, Josie Hintzke, Anne Gravel Sullivan, Graham Wills, Kayla Lemmon, Jason Dambach, Leigh Ann Mrotek, Mariah Quinn, Kirsten Abramson, Peter Kleinschmidt, Thomas B. Brazelton, Margaret A. Leaf, Heidi Twedt, David Kunstman, Brian Patterson, Frank Liao, Stacy Rasmussen, Elizabeth S. Burnside, Cherodeep Goswami, and Joel Gordon. A Pragmatic Randomized Controlled Trial of Ambient Artificial Intelligence to Improve Health Practitioner Well-Being. *NEJM AI*, 2(12): AIoa2500945, November 2025. doi: 10.1056/AIoa2500945. URL <https://ai.nejm.org/doi/10.1056/AIoa2500945>. Publisher: Massachusetts Medical Society.
- Baraa Al Jorf, Bartlomiej Piechowski-Jozwiak, and Farah E. Shamout. A data-centric perspective on designing AI foundation models for healthcare. *Frontiers in Digital Health*, Volume 8 - 2026, 2026.

- ISSN 2673-253X. doi: 10.3389/fdgth.2026.1738523. URL <https://www.frontiersin.org/journals/digital-health/articles/10.3389/fdgth.2026.1738523>.
- Seongsu Bae, Daeun Kyung, Jaehee Ryu, Eunbyeol Cho, Gyubok Lee, Sunjun Kweon, Jungwoo Oh, Lei Ji, Eric Chang, Tackeun Kim, and Edward Choi. EHRXQA: A Multi-Modal Question Answering Dataset for Electronic Health Records with Chest X-ray Images. *Advances in Neural Information Processing Systems*, 36:3867–3880, December 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/0c007ebef1d11fd48da6ce4f54687db6-Abstract-Datasets_and_Benchmarks.html.
- Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, Humen Zhong, Yuanzhi Zhu, Mingkun Yang, Zhaohai Li, Jianqiang Wan, Pengfei Wang, Wei Ding, Zheren Fu, Yiheng Xu, Jiabo Ye, Xi Zhang, Tianbao Xie, Zesen Cheng, Hang Zhang, Zhibo Yang, Haiyang Xu, and Junyang Lin. Qwen2.5-VL Technical Report, February 2025. URL <http://arxiv.org/abs/2502.13923>. arXiv:2502.13923 [cs].
- Brenton T. Bicknell, Danner Butler, Sydney Whalen, James Ricks, Cory J. Dixon, Abigail B. Clark, Olivia Spaedy, Adam Skelton, Neel Edupuganti, Lance Dzubinski, Hudson Tate, Garrett Dyess, Brenessa Lindeman, and Lisa Soleymani Lehmann. ChatGPT-4 Omni Performance in USMLE Disciplines and Clinical Skills: Comparative Analysis. *JMIR Medical Education*, 10(1):e63430, November 2024. doi: 10.2196/63430. URL <https://mededu.jmir.org/2024/1/e63430>.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language Models are Few-Shot Learners, May 2020. URL <https://arxiv.org/abs/2005.14165v4>.
- Queralt Miró Catalina, Aina Fuster-Casanovas, Josep Vidal-Alaball, Anna Escalé-Besa, Francesc X. Marin-Gomez, Joaquim Femenia, and Jordi Solé-Casals. Knowledge and perception of primary care healthcare professionals on the use of artificial intelligence as a healthcare tool. *DIGITAL HEALTH*, 9:20552076231180511, January 2023. ISSN 2055-2076. doi: 10.1177/20552076231180511. URL <https://doi.org/10.1177/20552076231180511>. Publisher: SAGE Publications Ltd.
- Mert Cemri, Melissa Z. Pan, Shuyi Yang, Lakshya A. Agrawal, Bhavya Chopra, Rishabh Tiwari, Kurt Keutzer, Aditya Parameswaran, Dan Klein, Kannan Ramchandran, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. Why Do Multi-Agent LLM Systems Fail?, March 2025. URL <https://arxiv.org/abs/2503.13657v3>.
- Emma Chen, Aman Kansal, Julie Chen, Boyang Tom Jin, Julia Reisler, David E Kim, and Pranav Rajpurkar. Multimodal Clinical Benchmark for Emergency Care (MC-BEC): A Comprehensive Benchmark for Evaluating Foundation Models in Emergency Medicine. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 45794–45811. Curran Associates, Inc., 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/8f61049e8fe5b9ed714860b951066f1e-Paper-Datasets_and_Benchmarks.pdf.
- Junde Chen, Yuxin Wen, Michael Pokojovy, Tzu-Liang (Bill) Tseng, Peter McCaffrey, Alexander Vo, Eric Walser, and Scott Moen. Multi-modal learning for inpatient length of stay prediction. *Computers in Biology and Medicine*, 171:108121, March 2024a. ISSN 0010-4825. doi: 10.1016/j.compbimed.2024.108121. URL <https://www.sciencedirect.com/science/article/pii/S001082524002051>.
- Junying Chen, Chi Gui, Ruyi Ouyang, Anningzhe Gao, Shunian Chen, Guiming Hardy Chen, Xidong Wang, Ruifei Zhang, Zhenyang Cai, Ke Ji, Guangjun Yu, Xiang Wan, and Benyou Wang. HuatuoGPT-Vision, Towards Injecting Medical Visual Knowledge into Multimodal LLMs at Scale, September 2024b. URL <http://arxiv.org/abs/2406.19280>. arXiv:2406.19280 [cs].

- Zhe Chen, Weiyun Wang, Yue Cao, Yangzhou Liu, Zhangwei Gao, Erfei Cui, Jinguo Zhu, Shenglong Ye, Hao Tian, Zhaoyang Liu, Lixin Gu, Xuehui Wang, Qingyun Li, Yiming Ren, Zixuan Chen, Jiapeng Luo, Jiahao Wang, Tan Jiang, Bo Wang, Conghui He, Botian Shi, Xingcheng Zhang, Han Lv, Yi Wang, Wenqi Shao, Pei Chu, Zhongying Tu, Tong He, Zhiyong Wu, Huipeng Deng, Jiaye Ge, Kai Chen, Kaipeng Zhang, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhai Wang. Expanding Performance Boundaries of Open-Source Multimodal Models with Model, Data, and Test-Time Scaling, September 2025. URL <http://arxiv.org/abs/2412.05271>. arXiv:2412.05271 [cs].
- Yilun Du, Shuang Li, Antonio Torralba, Joshua B. Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *ICML'24*, pages 11733–11763, Vienna, Austria, 2024. JMLR.org.
- Shaza Elsharief, Saeed Shurrab, Baraa Al Jorf, Leopoldo Julian Lechuga Lopez, Krzysztof J. Geras, and Farah E. Shamout. MedMod: Multimodal Benchmark for Medical Prediction Tasks with Electronic Health Records and Chest X-Ray Scans. In *Proceedings of the sixth Conference on Health, Inference, and Learning*, pages 781–803. PMLR, July 2025. URL <https://proceedings.mlr.press/v287/elsharief25a.html>.
- Mingyan Gao, Yanzi Li, Banruo Liu, Yifan Yu, Phillip Wang, Ching-Yu Lin, and Fan Lai. Single-agent or Multi-agent Systems? Why Not Both?, May 2025. URL <https://arxiv.org/abs/2505.18286v1>.
- Nasir Hayat, Krzysztof J. Geras, and Farah E. Shamout. MedFuse: Multi-modal fusion with clinical time-series data and chest X-ray images. In *Proceedings of the 7th Machine Learning for Healthcare Conference*, pages 479–503. PMLR, December 2022. URL <https://proceedings.mlr.press/v182/hayat22a.html>. ISSN: 2640-3498.
- Sirui Hong, Mingchen Zhuge, Jiaqi Chen, Xiawu Zheng, Yuheng Cheng, Ceyao Zhang, Jinlin Wang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang Zhou, Chenyu Ran, Lingfeng Xiao, Chenglin Wu, and Jürgen Schmidhuber. MetaGPT: Meta Programming for A Multi-Agent Collaborative Framework, August 2023. URL <https://arxiv.org/abs/2308.00352v7>.
- Yutai Hou, Hongyuan Dong, Xinghao Wang, Bohan Li, and Wanxiang Che. MetaPrompting: Learning to Learn Better Prompts. In Nicoletta Calzolari, Chu-Ren Huang, Hansaem Kim, James Pustejovsky, Leo Wanner, Key-Sun Choi, Pumo-Ryu, Hsin-Hsi Chen, Lucia Donatelli, Heng Ji, Sadao Kurohashi, Patrizia Paggio, Nianwen Xue, Seokhwan Kim, Younggyun Hahm, Zhong He, Tony Kyungil Lee, Enrico Santus, Francis Bond, and Seung-Hoon Na, editors, *Proceedings of the 29th International Conference on Computational Linguistics*, pages 3251–3262, Gyeongju, Republic of Korea, October 2022. International Committee on Computational Linguistics. URL <https://aclanthology.org/2022.coling-1.287/>.
- Kexin Huang, Jaan Altosaar, and Rajesh Ranganath. ClinicalBERT: Modeling Clinical Notes and Predicting Hospital Readmission, April 2019. URL <https://ui.adsabs.harvard.edu/abs/2019arXiv190405342H>. ADS Bibcode: 2019arXiv190405342H.
- Qiao Jin, Zhizheng Wang, Yifan Yang, Qingqing Zhu, Donald Wright, Thomas Huang, Nikhil Khandekar, Nicholas Wan, Xuguang Ai, W. John Wilbur, Zhe He, R. Andrew Taylor, Qingyu Chen, and Zhiyong Lu. AgentMD: Empowering language agents for risk prediction with large-scale clinical tool learning. *Nature Communications*, 16(1):9377, October 2025. ISSN 2041-1723. doi: 10.1038/s41467-025-64430-x. URL <https://www.nature.com/articles/s41467-025-64430-x>.
- Alistair Johnson, Tom Pollard, Steven Horng, Leo Anthony Celi, and Roger Mark. MIMIC-IV-Note: Deidentified free-text clinical notes, 2023a. URL <https://physionet.org/content/mimic-iv-note/2.2/>.
- Alistair E. W. Johnson, Tom J. Pollard, Seth J. Berkowitz, Nathaniel R. Greenbaum, Matthew P. Lungren, Chih-ying Deng, Roger G. Mark, and Steven Horng. MIMIC-CXR, a de-identified publicly available database of chest radiographs with free-text reports. *Scientific Data*, 6(1):317, December 2019. ISSN 2052-4463. doi: 10.1038/s41597-019-0322-0. URL <https://www.nature.com/a>

- rticles/s41597-019-0322-0. Publisher: Nature Publishing Group.
- Alistair E. W. Johnson, Lucas Bulgarelli, Lu Shen, Alvin Gayles, Ayad Shammout, Steven Horng, Tom J. Pollard, Sicheng Hao, Benjamin Moody, Brian Gow, Li-wei H. Lehman, Leo A. Celi, and Roger G. Mark. MIMIC-IV, a freely accessible electronic health record dataset. *Scientific Data*, 10(1):1, January 2023b. ISSN 2052-4463. doi: 10.1038/s41597-022-01899-x. URL <https://www.nature.com/articles/s41597-022-01899-x>.
- Baraa Al Jorf and Farah E. Shamout. Medpatch: Confidence-guided multi-stage fusion for multimodal clinical data. In Monica Agrawal, Kaivalya Deshpande, Matthew Engelhard, Shalmali Joshi, Shengpu Tang, and Iñigo Urteaga, editors, *Proceedings of the 10th Machine Learning for Healthcare Conference*, volume 298 of *Proceedings of Machine Learning Research*. PMLR, 15–16 Aug 2025. URL <https://proceedings.mlr.press/v298/jorf25a.html>.
- Lars Benedikt Kaesberg, Jonas Becker, Jan Philip Wahle, Terry Ruas, and Bela Gipp. Voting or Consensus? Decision-Making in Multi-Agent Debate. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Findings of the Association for Computational Linguistics: ACL 2025*, pages 11640–11671, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-256-5. doi: 10.18653/v1/2025.findings-acl.606. URL <https://aclanthology.org/2025.findings-acl.606/>.
- Beria Chingnabe Kalpelbe, Angel Gabriel Adaambiik, and Wei Peng. Vision Language Models in Medicine, February 2025. URL <http://arxiv.org/abs/2503.01863>. arXiv:2503.01863 [cs].
- Kaan Kara and Tuba Gunel. Clinical Risk Computation by Large Language Models Using Validated Risk Scores. *Journal of Medical Systems*, 49(1):121, September 2025. ISSN 1573-689X. doi: 10.1007/s10916-025-02261-5. URL <https://doi.org/10.1007/s10916-025-02261-5>.
- Firas Khader, Jakob Nikolas Kather, Gustav Müller-Franzes, Tianci Wang, Tianyu Han, Soroosh Tayebi Arasteh, Karim Hamesch, Keno Bressemer, Christoph Haarbuerger, Johannes Stegmaier, Christiane Kuhl, Sven Nebelung, and Daniel Truhn. Medical transformer for multimodal survival prediction in intensive care: integration of imaging and non-imaging data. *Scientific Reports*, 13(1):10666, July 2023. ISSN 2045-2322. doi: 10.1038/s41598-023-37835-1. URL <https://www.nature.com/articles/s41598-023-37835-1>.
- Yubin Kim, Chanwoo Park, Hyewon Jeong, Yik Siu Chan, Xuhai Xu, Daniel McDuff, Hyeonhoon Lee, Marzyeh Ghassemi, Cynthia Breazeal, and Hae Won Park. MDAgents: an adaptive collaboration of LLMs for medical decision-making. In *Proceedings of the 38th International Conference on Neural Information Processing Systems*, volume 37 of *NIPS '24*, pages 79410–79452, Red Hook, NY, USA, December 2024. Curran Associates Inc. ISBN 979-8-3313-1438-5.
- Yubin Kim, Ken Gu, Chanwoo Park, Chunjong Park, Samuel Schmidgall, A. Ali Heydari, Yao Yan, Zhihan Zhang, Yuchen Zhuang, Yun Liu, Mark Malhotra, Paul Pu Liang, Hae Won Park, Yuzhe Yang, Xuhai Xu, Yilun Du, Shwetak Patel, Tim Althoff, Daniel McDuff, and Xin Liu. Towards a Science of Scaling Agent Systems, December 2025. URL <https://arxiv.org/abs/2512.08296v3>.
- Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4):1234–1240, February 2020. ISSN 1367-4803. doi: 10.1093/bioinformatics/btz682. URL <https://doi.org/10.1093/bioinformatics/btz682>.
- Kwanhyung Lee, Soojeong Lee, Sangchul Hahn, Heejung Hyun, Edward Choi, Byungeun Ahn, and Joohyung Lee. Learning Missing Modal Electronic Health Records with Unified Multi-modal Data Embedding and Modality-Aware Attention. In *Proceedings of the 8th Machine Learning for Healthcare Conference*, pages 423–442. PMLR, December 2023. URL <https://proceedings.mlr.press/v219/lee23a.html>.
- Sujung Lee, Won Ik Cho, Youngrong Lee, Duck Ju Kim, Kyeng Hyun Nam, Sangmin Lee, Jungyo Suh, and Taehoon Ko. A prompt framework for enhancing LLM-based explainability of medical machine learning models: an intensive care unit application. *BMC Medical Informatics and Decision Making*, 25(1):430, November 2025. ISSN 1472-6947.

- doi: 10.1186/s12911-025-03239-6. URL <https://doi.org/10.1186/s12911-025-03239-6>.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Proceedings of the 34th International Conference on Neural Information Processing Systems, NIPS '20*, pages 9459–9474, Red Hook, NY, USA, 2020. Curran Associates Inc. ISBN 978-1-7138-2954-6. URL <https://dl.acm.org/doi/10.5555/3495724.3496517>.
- Chunyu Li, Cliff Wong, Sheng Zhang, Naoto Usuyama, Haotian Liu, Jianwei Yang, Tristan Naumann, Hoifung Poon, and Jianfeng Gao. LLaVA-Med: Training a Large Language-and-Vision Assistant for Biomedicine in One Day, June 2023. URL <http://arxiv.org/abs/2306.00890>. arXiv:2306.00890 [cs].
- Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Shuming Shi, and Zhaopeng Tu. Encouraging Divergent Thinking in Large Language Models through Multi-Agent Debate. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen, editors, *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 17889–17904, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.992. URL <https://aclanthology.org/2024.emnlp-main.992/>.
- Shiyin Lu, Yang Li, Yu Xia, Yuwei Hu, Shanshan Zhao, Yanqing Ma, Zhichao Wei, Yinglun Li, Lunhao Duan, Jianshan Zhao, Yuxuan Han, Haijun Li, Wanying Chen, Junke Tang, Chengkun Hou, Zhixing Du, Tianli Zhou, Wenjie Zhang, Huping Ding, Jiahe Li, Wen Li, Gui Hu, Yiliang Gu, Siran Yang, Jiamang Wang, Hailong Sun, Yibo Wang, Hui Sun, Jinlong Huang, Yuping He, Shengze Shi, Weihong Zhang, Guodong Zheng, Junpeng Jiang, Sensen Gao, Yi-Feng Wu, Sijia Chen, Yuhui Chen, Qing-Guo Chen, Zhao Xu, Weihua Luo, and Kaifu Zhang. Ovis2.5 technical report. *arXiv:2508.11737*, 2025.
- Paul J. Lukac, William Turner, Sitaram Vangala, Aaron T. Chin, Joshua Khalili, Ya-Chen Tina Shih, Catherine Sarkisian, Eric M. Cheng, and John N. Mafi. Ambient AI Scribes in Clinical Practice: A Randomized Trial. *NEJM AI*, 2(12):AIoa2501000, November 2025. doi: 10.1056/AIoa2501000. URL <https://ai.nejm.org/doi/full/10.1056/AIoa2501000>. Publisher: Massachusetts Medical Society.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhunoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. Self-Refine: Iterative Refinement with Self-Feedback, May 2023. URL <http://arxiv.org/abs/2303.17651>. arXiv:2303.17651 [cs].
- Harsha Nori, Yin Tat Lee, Sheng Zhang, Dean Carignan, Richard Edgar, Nicolo Fusi, Nicholas King, Jonathan Larson, Yuanzhi Li, Weishung Liu, Renqian Luo, Scott Mayer McKinney, Robert Osazuwa Ness, Hoifung Poon, Tao Qin, Naoto Usuyama, Chris White, and Eric Horvitz. Can Generalist Foundation Models Outcompete Special-Purpose Tuning? Case Study in Medicine, November 2023. URL <http://arxiv.org/abs/2311.16452>. arXiv:2311.16452 [cs].
- Andrew Sellergren, Sahar Kazemzadeh, Tiam Jaroensri, Atilla Kiraly, Madeleine Traverse, Timo Kohlberger, Shawn Xu, Fayaz Jamil, Cian Hughes, Charles Lau, Justin Chen, Fereshteh Mahvar, Liron Yatziv, Tiffany Chen, Bram Sterling, Stefanie Anna Baby, Susanna Maria Baby, Jeremy Lai, Samuel Schmidgall, Lu Yang, Kejia Chen, Per Bjornsson, Shashir Reddy, Ryan Brush, Kenneth Philbrick, Mercy Asiedu, Ines Mezerreg, Howard Hu, Howard Yang, Richa Tiwari, Sunny Jansen, Preeti Singh, Yun Liu, Shekoofeh Azizi, Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, Tatiana Matejovicova, Alexandre Ramé, Morgane Riviere, Louis Rouillard, Thomas Mesnard, Geoffrey Cideron, Jean-bastien Grill, Sabela Ramos, Edouard Yvinec, Michelle Casbon, Elena Buchatskaya, Jean-Baptiste Alayrac, Dmitry Lepikhin, Vlad Feinberg, Sebastian Borgeaud, Alek Andreev, Cassidy Hardin, Robert Dadashi, Léonard Hussenot, Armand Joulin, Olivier Bachem, Yossi Matias, Katherine Chou, Avinatan Hassidim, Kavi Goel, Clement Farabet, Joelle Barral, Tris Warkentin, Jonathon Shlens, David Fleet, Victor Cotruta,

- Omar Sanseviero, Gus Martins, Phoebe Kirk, Anand Rao, Shravya Shetty, David F. Steiner, Can Kirmizibayrak, Rory Pilgrim, Daniel Golden, and Lin Yang. MedGemma Technical Report, July 2025. URL <http://arxiv.org/abs/2507.05201>. arXiv:2507.05201 [cs].
- Ali Shuaib. Transforming Healthcare with AI: Promises, Pitfalls, and Pathways Forward. *International Journal of General Medicine*, 17:1765–1771, May 2024. doi: 10.2147/IJGM.S449598. URL <https://www.dovepress.com/transforming-healthcare-with-ai-promises-pitfalls-and-pathways-forward-peer-reviewed-fulltext-article-IJGM>. Publisher: Dove Press.
- Karan Singhal, Shekoofeh Azizi, Tao Tu, S. Sara Mahdavi, Jason Wei, Hyung Won Chung, Nathan Scales, Ajay Tanwani, Heather Cole-Lewis, Stephen Pfohl, Perry Payne, Martin Seneviratne, Paul Gamble, Chris Kelly, Abubakr Babiker, Nathanael Schärli, Aakanksha Chowdhery, Philip Mansfield, Dina Demner-Fushman, Blaise Agüera y Arcas, Dale Webster, Greg S. Corrado, Yossi Matias, Katherine Chou, Juraj Gottweis, Nenad Tomasev, Yun Liu, Alvin Rajkomar, Joelle Barral, Christopher Semturs, Alan Karthikesalingam, and Vivek Natarajan. Large language models encode clinical knowledge. *Nature*, 620 (7972):172–180, August 2023. ISSN 1476-4687. doi: 10.1038/s41586-023-06291-2. URL <https://www.nature.com/articles/s41586-023-06291-2>.
- Karan Singhal, Tao Tu, Juraj Gottweis, Rory Sayres, Ellery Wulczyn, Mohamed Amin, Le Hou, Kevin Clark, Stephen R. Pfohl, Heather Cole-Lewis, Darlene Neal, Qazi Mamunur Rashid, Mike Schaekermann, Amy Wang, Dev Dash, Jonathan H. Chen, Nigam H. Shah, Sami Lachgar, Philip Andrew Mansfield, Sushant Prakash, Bradley Green, Ewa Dominowska, Blaise Agüera y Arcas, Nenad Tomašev, Yun Liu, Renee Wong, Christopher Semturs, S. Sara Mahdavi, Joelle K. Barral, Dale R. Webster, Greg S. Corrado, Yossi Matias, Shekoofeh Azizi, Alan Karthikesalingam, and Vivek Natarajan. Toward expert-level medical question answering with large language models. *Nature Medicine*, 31(3):943–950, March 2025. ISSN 1546-170X. doi: 10.1038/s41591-024-03423-7. URL <https://www.nature.com/articles/s41591-024-03423-7>. Publisher: Nature Publishing Group.
- Ming Tai-Seale, Sally L. Baxter, Florin Vaida, Amanda Walker, Amy M. Sitapati, Chad Osborne, Joseph Diaz, Nimit Desai, Sophie Webb, Gregory Polston, Teresa Helsten, Erin Gross, Jessica Thackaberry, Ammar Mandvi, Dustin Lillie, Steve Li, Geneen Gin, Suraj Achar, Heather Hofflich, Christopher Sharp, Marlene Millen, and Christopher A. Longhurst. AI-Generated Draft Replies Integrated Into Health Records and Physicians’ Electronic Communication. *JAMA Network Open*, 7 (4):e246565, April 2024. ISSN 2574-3805. doi: 10.1001/jamanetworkopen.2024.6565. URL <https://doi.org/10.1001/jamanetworkopen.2024.6565>.
- Mingtian Tan, Mike A. Merrill, Vinayak Gupta, Tim Althoff, and Thomas Hartvigsen. Are language models actually useful for time series forecasting? In *Proceedings of the 38th International Conference on Neural Information Processing Systems*, volume 37 of *NIPS ’24*, pages 60162–60191, Red Hook, NY, USA, December 2024. Curran Associates Inc. ISBN 979-8-3313-1438-5.
- Xiangru Tang, Anni Zou, Zhuosheng Zhang, Ziming Li, Yilun Zhao, Xingyao Zhang, Arman Cohan, and Mark Gerstein. MedAgents: Large Language Models as Collaborators for Zero-shot Medical Reasoning. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Findings of the Association for Computational Linguistics: ACL 2024*, pages 599–621, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.33. URL <https://aclanthology.org/2024.findings-acl.33/>.
- Warren J. von Eschenbach. Transparency and the Black Box Problem: Why We Do Not Trust AI. *Philosophy & Technology*, 34(4):1607–1622, December 2021. ISSN 2210-5441. doi: 10.1007/s13347-021-00477-0. URL <https://doi.org/10.1007/s13347-021-00477-0>.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-Consistency Improves Chain of Thought Reasoning in Language Models, March 2022. URL <https://arxiv.org/abs/2203.11171v4>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting

elicits reasoning in large language models. In *Proceedings of the 36th International Conference on Neural Information Processing Systems, NIPS '22*, pages 24824–24837, Red Hook, NY, USA, November 2022. Curran Associates Inc. ISBN 978-1-7138-7108-8.

Yingxuan Yang, Huacan Chai, Shuai Shao, Yuanyi Song, Siyuan Qi, Renting Rui, and Weinan Zhang. AgentNet: Decentralized Evolutionary Coordination for LLM-based Multi-Agent Systems, April 2025. URL <https://arxiv.org/abs/2504.00587v2>.

Sihang Zeng, Yujuan Fu, Sitong Zhou, Zixuan Yu, Lucas Jing Liu, Jun Wen, Matthew Thompson, Ruth Etzioni, and Meliha Yetisgen. Traj-coa: Patient trajectory modeling via chain-of-agents for lung cancer risk prediction, 2025. URL <https://arxiv.org/abs/2510.10454>. Accepted by NeurIPS 2025 GenAI4Health Workshop.

Yinghao Zhu, Ziyi He, Haoran Hu, Xiaochen Zheng, Xichen Zhang, Zixiang Wang, Junyi Gao, Liantao Ma, and Lequan Yu. MedAgentBoard: Benchmarking Multi-Agent Collaboration with Conventional Methods for Diverse Medical Tasks. October 2025. URL <https://openreview.net/forum?id=BPpG4qQaNj>.

Appendix A. Algorithms

This section shows all the algorithms used for each agentic setup. Single agents follow algorithm [A1](#) in the unimodal setup and algorithm [A2](#) in the multimodal setup. Multi-agent frameworks follow algorithm [A3](#).

Algorithm A1: Single Agent Unimodal (PS Only)

Input: Patient Summary \mathbf{x}_{ps} , System Prompt \mathcal{P}_{task}

Output: Prediction \hat{y}

$Context \leftarrow \text{concat}(\mathcal{P}_{task}, \mathbf{x}_{ps})$

$\hat{y} \leftarrow \mathcal{A}(Context)$

return \hat{y}

Algorithm A2: Single Agent Multimodal

Input: Available modalities $\mathcal{X}_{available}$,

System Prompt \mathcal{P}_{task}

Output: Prediction \hat{y}

$Context \leftarrow \mathcal{P}_{task}$

foreach $\mathbf{x}_m \in \mathcal{X}_{available}$ **do**

 | $Context \leftarrow \text{concat}(Context, \text{"Modality } m: ", \mathbf{x}_m)$

end

$\hat{y} \leftarrow \mathcal{A}(Context)$

return \hat{y}

Algorithm A3: Multi-Agent Multimodal

Input: Available modalities $\mathcal{X}_{available}$,

System Prompt \mathcal{P}_{task}

Output: Prediction \hat{y}

$S \leftarrow 0$

foreach $\mathbf{x}_m \in \mathcal{X}_{available}$ **do**

 | $Context_m \leftarrow \text{concat}(\mathcal{P}_{task}, \mathbf{x}_m)$

 | $p_m \leftarrow \mathcal{A}_m(Context_m)$

 | $S \leftarrow S + p_m$

end

$\bar{p} \leftarrow S/|\mathcal{X}_{available}|$

$\hat{y} \leftarrow \mathbb{I}(\bar{p} > 0.5)$

return \hat{y}

Appendix B. Additional Results

B.1. Intern Multimodal

Table B1 details the performance of the InternVL2.5 model. Note that this architecture supports three modalities (PS, CXR, RR), excluding EHR. We report the Area Under the Receiver Operating Characteristic (AUROC), Area Under the Precision-Recall Curve (AUPRC), and Expected Calibration Error (ECE) for both In-Hospital Mortality and Length of Stay (> 7 Days) prediction tasks. Figure B2 illustrates the shift in ECE when transitioning from the unimodal to the multimodal setting within single-agent frameworks. Consistent with our findings, the chart highlights that adding modalities generally improves calibration for robust backbones, though this benefit varies across model architectures.

Table B1: **InternVL2.5 Multimodal Results.** Results for InternVL2.5 are reported separately as this model supports only 3 modalities (missing Timeseries). Best overall performance in each column is bolded.

Backbone	Arch.	Method	In-Hospital Mortality			Length of Stay (>7 Days)		
			AUROC	AUPRC	ECE	AUROC	AUPRC	ECE
Intern	Single	Zero-shot	0.700 (0.680 - 0.723)	0.247 (0.223 - 0.281)	0.082	0.721 (0.703 - 0.738)	0.358 (0.331 - 0.387)	0.354
		Few-shot	0.745 (0.724 - 0.766)	0.282 (0.252 - 0.315)	0.121	0.677 (0.658 - 0.695)	0.302 (0.280 - 0.328)	0.232
		CoT	0.683 (0.662 - 0.705)	0.216 (0.194 - 0.242)	0.094	0.674 (0.655 - 0.693)	0.319 (0.293 - 0.348)	0.445
		CoT-SC	0.727 (0.707-0.748)	0.270 (0.707 - 0.748)	0.032	0.690 (0.670 - 0.709)	0.345 (0.319 - 0.377)	0.456
	Multi	Majority Vote	0.709 (0.688 - 0.731)	0.249 (0.220 - 0.282)	0.106	0.679 (0.658 - 0.697)	0.306 (0.281 - 0.335)	0.056
		Debate (Uni)	0.627 (0.603 - 0.651)	0.181 (0.161 - 0.207)	0.053	0.641 (0.620 - 0.660)	0.285 (0.259 - 0.313)	0.006
		Meta-Prompt	0.554 (0.529 - 0.580)	0.142 (0.128 - 0.159)	0.223	0.550 (0.528 - 0.571)	0.230 (0.212 - 0.251)	0.374

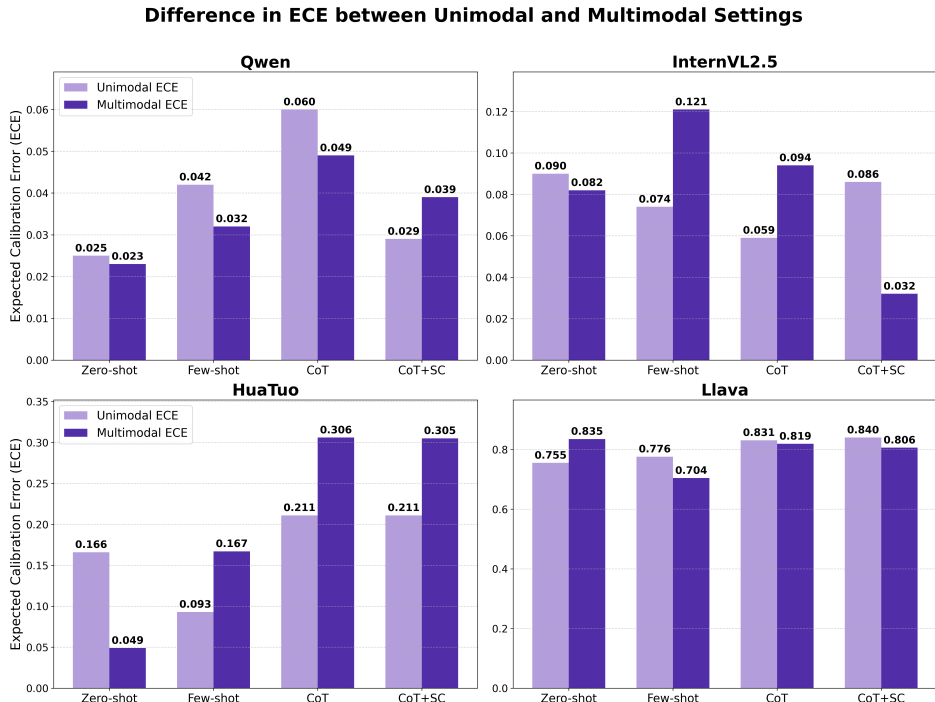


Figure B2: Bar graphs showing the difference in ECE between the unimodal and multimodal settings for the single agent frameworks across the four backbones.

Appendix C. Additional Ablations

C.1. Modality Addition Ablations

In this section, we provide visual support for the calibration and performance trends observed in our ablation studies. Figure C1 offers a holistic comparison between the effect of adding more modalities to the single agent and multi-agent setups. The improvement of the single agent configuration as more modalities are added compared to the drop in majority vote reinforces our conclusion that unified context processing synthesizes heterogeneous clinical data more effectively than decentralized mechanisms.

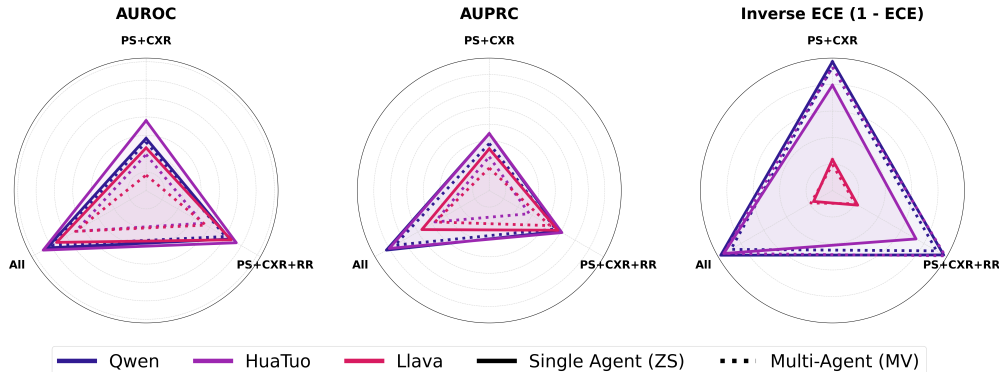


Figure C1: Radar charts showing how the single agent setup benefits more from the multimodality than the multiagent setup. Each circle corresponds to a metric. The dotted lines correspond to majority vote and the solid ones to zero-shot. Note the the chart shows the inverse of ECE for consistency with the other metrics.

C.2. Weighted Majority Vote

We examine the effect of weighting each modality agent by its unimodal AUROC performance in the Majority vote backbone. The results are shown in table C2

Table C2: **Performance Results.** Comparison of Unimodal baselines and Multimodal Multi-Agent voting methods. Best overall performance in each column is bolded.

Backbone	Arch.	Method	Performance	
			AUROC	AUPRC
Qwen	Unimodal	PS	0.666	0.232
		RR	0.684	0.222
		EHR	0.662	0.230
		CXR	0.538	0.141
	Multimodal Multi-Agent	Majority Vote	0.748	0.315
		Weighted Majority Vote	0.750	0.318

C.3. Sycophancy Analysis

We examine the reasoning traces from the debate between unimodal agents baseline and summarize our findings in Table C3. Notably, the agents seem to exhibit varying levels of sycophancy and echo chamber agreement depending on the backbone used.

Table C3: Debate traces analysis and echo chamber behavior per backbone.

Backbone	Pattern	Evidence
Qwen	Premature agreement	Consensus in round 1 for all 4,925 cases.
Huatuo	Weak echo chamber	Some later consensuses aligned with the initial majority.
MedGemma	Strong echo chamber	75/76 initial disagreements later converged. All followed the round-1 majority.
LLaVaMed	Unstable communication	Disagreements rarely resolved into stable consensus.

C.4. Debate Between Multimodal Agents

In Table C4, we examine the performance of debate between multimodal agents. Each agent in this setup receives the full multimodal context and debates with the other multimodal agents until consensus.

Table C4: **Debate Multimodal Results.** Comparison of in-hospital mortality and length of stay prediction performance for HuaTuo and Qwen using the debate architecture with multimodal agents. Best overall performance in each column is bolded.

Backbone	In-Hospital Mortality		Length of Stay (>7 Days)	
	AUROC	AUPRC	AUROC	AUPRC
HuaTuo	0.684 (0.663 - 0.707)	0.235 (0.210 - 0.266)	0.592 (0.572 - 0.613)	0.249 (0.231 - 0.272)
Qwen	0.707 (0.686 - 0.728)	0.243 (0.215 - 0.275)	0.561 (0.541 - 0.581)	0.231 (0.213 - 0.252)

C.5. Note-based Task

In Table C5, we test the performance of simple single agent and multi-agent setups on a note-based chronic kidney disease detection task. We extract the labels from MIMIC IV.

Table C5: **Multimodal Results.** Comparison of Unimodal, Multimodal, and Multi-Agent architectures for Chronic Kidney Disease Detection. Best overall performance in each column is bolded.

Backbone	Arch.	Method	Chronic Kidney Disease Detection	
			AUROC	AUPRC
Supervised	Unimodal	BioBERT	0.878	0.705
	Multimodal	MedPatch	0.807	0.532
Qwen	Unimodal Single Agent	Zero-shot	0.968	0.924
	Multimodal Multi-Agent	Majority Vote	0.959	0.897

C.6. MedGemma

Table C6 shows the results on the in-hospital mortality task using the MedGemma 4B backbone (Søllergren et al., 2025). We note that performance does not diverge much from the larger backbones.

Table C6: **Multimodal Results.** Comparison of Unimodal Single-Agent, Multimodal Single-Agent, and Multimodal Multi-Agent frameworks on the mortality task using the MedGemma backbone. Best overall performance in each column is bolded.

Backbone	Arch.	Method	In-Hospital Mortality	
			AUROC	AUPRC
MedGemma	Unimodal Single Agent	Zero-shot	0.626 (0.602 - 0.649)	0.179 (0.160 - 0.201)
		Few-shot	0.691 (0.207 - 0.714)	0.232 (0.207 - 0.263)
		CoT	0.624 (0.600 - 0.647)	0.172 (0.154 - 0.190)
		Self-Refine	0.628 (0.603 - 0.652)	0.183 (0.163 - 0.206)
	Multimodal Single Agent	Zero-shot	0.735 (0.714 - 0.755)	0.250 (0.225 - 0.282)
		Few-shot	0.744 (0.723 - 0.765)	0.280 (0.250 - 0.314)
	Multimodal Multi-Agent	Majority Vote	0.687 (0.664 - 0.708)	0.226 (0.201 - 0.256)
		Debate	0.594 (0.569 - 0.619)	0.172 (0.151 - 0.197)

C.7. EHR Serialization

Table C7 details an ablation on EHR serialization methods. We tested our standard log baseline against two reduced-context approaches: the summary method (aggregating extreme and average baseline values) and the delta method (representing only the net clinical trajectory). Best overall performance in each column is bolded.

Table C7: **Performance Results.** Comparison of Log, Summary, and Delta methods.

Method	AUROC	AUPRC
Original Log	0.756	0.330
Summary	0.760	0.301
Delta	0.760	0.303